

Simple Linear Regression (3.1.1)

```
In [46]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [47]: data = pd.read_csv('./Salary_Data.csv')
data.head()
```

```
Out[47]:
```

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

```
In [48]: x = data.loc[:, ['YearsExperience']]
y = data.loc[:, ['Salary']]
```

```
In [49]: print(x.head())
print(y.head())
```

```
YearsExperience
0      1.1
1      1.3
2      1.5
3      2.0
4      2.2

Salary
0    39343
1    46205
2    37731
3    43525
4    39891
```

```
In [8]: divide = np.random.rand(len(data))<=0.8
divide
```

```
Out[8]: array([ True, False, False,  True,  True,  True,  True,  True, False,
        True, False,  True, False,  True,  True,  True,  True, False,
        True,  True,  True,  True,  True,  True,  True,  True,  True,
        True,  True, False])
```

```
In [11]: training_set = data[divide]
testing_set = data[~divide]
```

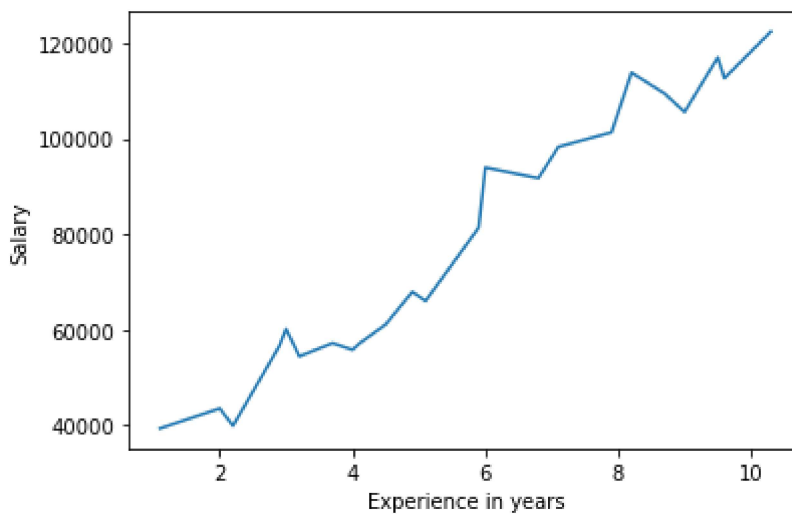
```
In [12]: print(testing_set)
```

```
YearsExperience  Salary
1              1.3    46205
```

2	1.5	37731
8	3.2	64445
10	3.9	63218
12	4.0	56957
17	5.3	83088
29	10.5	121872

```
In [13]: x_train = training_set.iloc[:,0]
x_test = testing_set.iloc[:, 0]
y_train = training_set.iloc[:,1]
y_test = testing_set.iloc[:, 1]
```

```
In [14]: plt.plot(x_train, y_train)
plt.xlabel("Experience in years")
plt.ylabel("Salary")
plt.show()
```



```
In [25]: def linear_regression(x, y, learning_rate, iterations, theta0, theta1):
n = float(len(x))
for i in range(iterations):
    prec_y = theta0 + theta1*x
    theta1 = theta1 - (-2/n)*np.sum(x * (y - prec_y))*learning_rate
    theta0 = theta0 - (-2/n)*np.sum(y - prec_y)*learning_rate
return (theta0, theta1)
```

```
In [26]: print(linear_regression(x_train, y_train, 0.0001, 10000, 0, 0))
```

(8402.21922139043, 11958.428699413906)

```
In [27]: print(linear_regression(x_train, y_train, 0.0001, 10000, 0, 0)[0])
```


8402.21922139043

```
In [41]: (theta0, theta1) = linear_regression(x_train, y_train, 0.0002, 30000, 0, 0)
```

```
In [42]: print(f"Slope: {theta1}\nIntercept: {theta0}")
```

Slope: 10136.852526872284
Intercept: 20938.593429415432

```
In [43]: prec_y = theta0 + theta1*x_train
plt.scatter(x_train, y_train)
plt.xlabel("Experience in years")
plt.ylabel("Salary")
plt.plot([min(x_train), max(x_train)], [min(prec_y), max(prec_y)])
plt.show()
```



```
In [44]: model_prec = theta0 + theta1*x_test
model_prec #testing data output from model parameters
```

```
Out[44]: 1      34116.501714
2      36143.872220
8      53376.521515
10     60472.318284
12     61486.003537
17     74663.911822
29     127375.544962
Name: YearsExperience, dtype: float64
```

```
In [45]: n = len(x_test)
meanSquare = (1/n)*np.sum(np.square(y_test-model_prec))
cost = np.sqrt(meanSquare)
print(cost)
```

7563.699881257977

Multiple Linear Regression (3.1.2)

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: data = pd.read_csv('./housingPrices.csv', header=None, delim_whitespace=True)
n = len(data)
data.shape
```

```
Out[2]: (506, 14)
```

```
In [3]: train_data = data.sample(frac = 0.8, random_state = 25)
test_data = data.drop(train_data.index)
print(train_data.shape)
```

```
(405, 14)
```

```
In [4]: train_data = train_data.values
        test_data = test_data.values
        print(test_data.shape)
```

```
(101, 14)
```

```
In [5]: train_y = train_data[:, -1].reshape(train_data.shape[0], 1) # all rows and last column
        train_x = train_data[:, :-1] # all rows and remaining columns
```

```
In [6]: test_y = test_data[:, -1].reshape(test_data.shape[0], 1)
        test_x = test_data[:, :-1]
```

```
In [7]: train_x = np.c_[np.ones(train_x.shape[0]), train_x]
        test_x = np.c_[np.ones(test_x.shape[0]), test_x]
        print(train_x.shape)
        print(test_x.shape)
```

```
(405, 14)
```

```
(101, 14)
```

```
In [8]: def gradient_descent(x, y, learning_rate, iterations):
        m = y.size
        theta = np.zeros((x.shape[1], 1))
        # cost_list = []
        for i in range(iterations):
            prec_y = np.dot(x, theta)
            cost = (1/(2*m))*np.sum(np.square(prec_y - y))
            theta = theta - learning_rate*(1/m)*np.dot(x.T, prec_y-y)
        # cost_list.append(cost)
        if(i % (iterations/10) == 0):
            print(f"Cost: {cost}")
        return theta
```

```
In [10]: theta = gradient_descent(train_x, train_y, 0.000005, 500000)
```

```
Cost: 298.84829629629627
Cost: 19.84771688334405
Cost: 17.678868134719522
Cost: 16.140328865301385
Cost: 15.024895429668973
Cost: 14.214835703851927
Cost: 13.625728473057716
Cost: 13.196660559812997
Cost: 12.883608654558083
Cost: 12.654717844389664
```

```
In [11]: yTestPrec = np.dot(test_x, theta)
```

```
In [12]: m = yTestPrec.shape[0]
        error = (1/m)*np.sum(np.square(yTestPrec - test_y))
        error = np.sqrt(error)
        print(error)
```

5.3297955443715415